

Large Scale Kernel methods

- ▶ Speaker: Issam Laradji
- ▶ MLRG

Outline

- ▶ Introduction to kernel methods
- ▶ Low rank approximation
 - ▶ Nystrom approximation
- ▶ Random Fourier Features
 - ▶ Random Kitchen Sink
 - ▶ Fastfood

Kernel

- ▶ Dataset $X \in R^{N \times D}$
 - ▶ N samples, D features
- ▶ A similarity function that takes two input vectors and spits out their similarity

$$K(X_i, X_j) = \phi(X_i) \cdot \phi(X_j)$$

- ▶ $\phi(X_i)$ is the feature representation of X_i in the higher dimensional space (possibly infinite)
 - ▶ Need not be explicitly computed
- ▶ Why kernels ?
 - ▶ $\phi(X_i) \cdot \phi(X_j)$ may have high dimensional complexity
 - ▶ $\phi(X_i) \cdot \phi(X_j)$ might be impossible to compute.
 - ▶ Gaussian Kernel: $\phi(X_j)$ projects X_j to infinite dimensions

Kernel

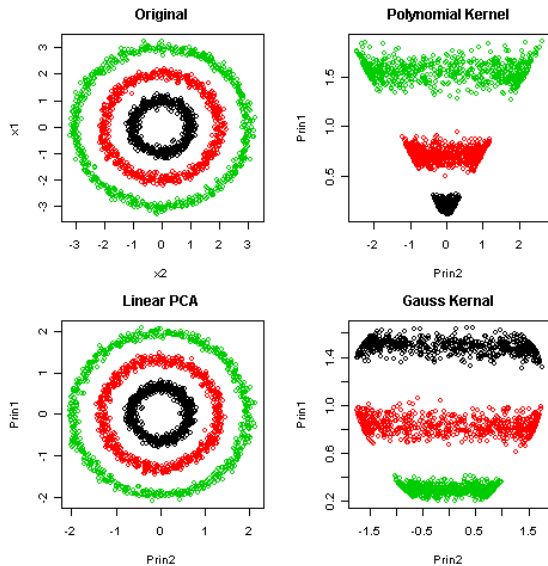


Figure 1:

Example

- ▶ Gaussian Kernel:

- ▶ As, $K(X_i, X_j) = e^{-\frac{||X_i - X_j||^2}{2\sigma^2}}$
- ▶ Also as, $K(X_i, X_j) = e^{-\frac{X_i \cdot X_j}{\sigma^2}}$
- ▶ which is $K(X_i, X_j) = \sum_{n=0}^{\infty} \frac{(X_i \cdot X_j)^n}{\sigma^n n!}$

- ▶ Gaussian kernel is a combination of all polynomial kernels of degrees $n \geq 0$
- ▶ It's an infinite power series which converges

Kernel-based algorithms

- ▶ Dataset $X \in R^{N \times D}$, $y \in R^N$
 - ▶ N samples, D features
- ▶ SVMs, Kernel Ridge Regression, Gaussian Process Regression
- ▶ A Kernel matrix is a Gram matrix $K \in R^{N \times N}$
- ▶ Linear function expansion is

$$f(x) = \sum_i w_i K(X_i, X)$$

- ▶ Therefore, number of basis functions increases linearly in the number of observations
- ▶ Kernel Ridge Regression

$$w = [K + \lambda I]^{-1} z y$$

- ▶ where $z_i = K(X_i, X)$
- ▶ Computational cost for large-scale problems
 - ▶ $\Omega(N^2)$ space.
 - ▶ $O(N^3)$ time for matrix inversion or SVD

Example

- ▶ Inverting a large matrix with $N = 18\text{M}$
- ▶ $K \approx 1300$ TB
- ▶ $320,000 \times 4$ GB RAM machines

Two main strategies

- ▶ Sampling-based low-rank approximation:
 - ▶ Compute and store only $T \ll N$ columns of K
 - ▶ Column-sampling using Nystrom method
- ▶ Explicit feature expansion (Fast feature extraction):
 - ▶ Fourier random features
 - ▶ Random Kitchen Sink, Fastfood

Original Nystrom Approximation

- ▶ Positive definite symmetric matrix $K \in R^{N \times N}$
- ▶ Computing the inverse - SVD on very large datasets can become prohibitive quickly for large N $O(N^3)$
- ▶ The Nystrom method is an efficient technique for the eigenvalue decomposition of large kernel matrices

Algorithm

- ▶ Input : Gram matrix $K \in \mathbb{R}^{N \times N}$
 - ▶ Result : $\hat{K} \in \mathbb{R}^{N \times N}$
1. Pick T columns of K in i.i.d trials, uniformly with replacement
 2. Let I be the set of indices of the sampled columns
 3. Let C be the $N \times T$ matrix containing the sampled columns
 4. Let H be the $T \times T$ submatrix of K whose entries are K_{ij} ,
 $i \in I, j \in I$
 5. Return $\hat{K} = CH^{-1}C$
- ▶ Computes an approximation $\hat{K} \approx K$
 - ▶ Computing SVD of H
 $O(T^3)$
 - ▶ Computing \hat{K}
 $O(N \cdot T \cdot k)$
 - ▶ where k is the best rank- k approximation of K
 - ▶ $k \leq T$

Nyström Woodbury Approximation

- ▶ Kernel ridge regression

$$w = [K + \lambda I]^{-1} z y$$

- ▶ Computing $\underbrace{[K + \lambda I]^{-1}}_{O(N^3)}$

- ▶ Matrix inversion lemma

$$\begin{aligned}(K + \lambda I)^{-1} &\approx (\hat{K} + \lambda I)^{-1} \\&= (CH^{-1}C^T + \lambda I)^{-1} \\&= \frac{1}{\lambda}(I - C(\lambda I + H^{-1}C^T C)^{-1}H^{-1}C^T) \\&= \frac{1}{\lambda}(I - C\underbrace{(\lambda I + H^{-1}C^T C)^{-1}}_{O(T^3)}H^{-1}C^T)\end{aligned}$$

- ▶ Inverting a $T \times T$ matrix instead

(Williams & Seeger, 2000)

Approximation Error

Drineas & Mahoney (2005)

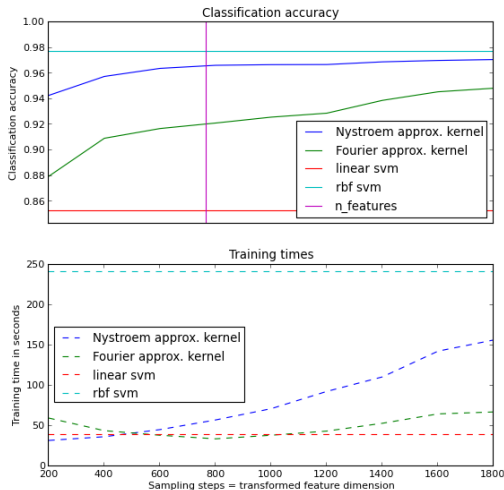
Theorem 1:

- ▶ Let \hat{K} be the best rank- k approximation of K
- ▶ Let $b = 1 + \sqrt{8 \log(1/\delta)}$
- ▶ Let ϵ be a small value larger than 0
- ▶ If $T \geq 4b^2/\epsilon^2$, then with probability at least $1 - \delta$

$$E[\|K - \hat{K}\|_F] \leq \|K - \hat{K}\|_F + \epsilon \sum_{i=1}^N K_{ii}^2$$

Example

- ▶ MNIST Dataset
 - ▶ 20000 training examples and 10000 test examples.



Applications

► Examples

- Spectral Clustering (Fowlkes et al., 2004).
- Kernel Ridge Regression (Cortes, MM, and Talwalkar, AISTATS 2010).
- Support Vector Machines (Fine and Scheinberg, 2001).
- Kernel Logistic Regression (Karsmarker et al., 2007).
- Manifold Learning (Kumar and Talwalkar, 2008)

Extensions on Nystrom

- ▶ We discussed uniform sampling of columns
- ▶ Other sampling methods include:
 - ▶ ℓ_2 norm of the columns
 - ▶ Non-uniformly \propto diagonal elements K_{ii}
- ▶ Empirical results:
 - ▶ uniform sampling without replacement: best results and fastest for real-world datasets (Kumar, MM, and Talwalkar, 2009).
- ▶ Found to work great for sparse optimization

$$\min_w \frac{1}{2n} \|y - Kw\|_2^2 + \lambda \|w\|_1$$

2. Fourier Random Features

- ▶ Linear decision surface over the kernelized form,

$$f(x) = \sum_{i=1}^N w_i K(X_i, X)$$

- ▶ The Gram matrix K can be too large (10,000, 10,000)




Figure 4:

2. Fourier Random Features

- ▶ Linear decision surface over the kernelized form,

$$f(x) = \sum_{i=1}^N w_i K(X_i, X)$$

- ▶ $K(X_i, X_j) = \phi(X_i) \cdot \phi(X_j)$
- ▶ Main Idea: Compute an approximated representation of $\phi(X_i)$
 - ▶ has $O(\log N)$ dimensional space
- ▶ Instead of computing $K \in R^{N \times N}$, compute $\phi(X) \in R^{N \times \log(N)}$
- ▶ For 10,000 images, $\log_{10}(10^4) = 4$
- ▶ Exact computation of $\phi(X)$




Figure 5:

2. Fourier Random Features

- Compute an approximation $\phi_z(X) \approx \phi(X)$ explicitly




Figure 6:

- $K(X_i, X_j) = \phi(X_i) \cdot \phi(X_j)$
- Compute $\phi_z(X_i) \cdot \phi_z(X_j)$ such that

$$E_z[\phi_z(X_i)\phi_z(X_j)] = K(X_i, X_j)$$

2. Fourier Random Features

- ▶ Let K be a translation invariant kernel, that is,

$$K(X_i, X_j) = K(X_i - X_j, 0)$$

- ▶ such as the gaussian kernel
- ▶ Bochner's Theorem: A kernel $K(X_i - X_j)$ is PSD iff $K(X_i - X_j)$ is the Fourier transform of a non-negative measure $p(z)$

$$K(X_i - X_j, 0) = \int_{\mathbb{R}} e^{-iz(X_i - X_j)} p(z) dz$$

- ▶ where $p(z)$ is a positive finite measure which is when scaled is a proper probability distribution.

2. Fourier Random Features

- Therefore,

$$\begin{aligned}k(X_i, X_j) &= k(X_i - X_j, 0) \\&= \int_{R^d} p(z) e^{z(X_i - X_j)} dz \\&\approx \frac{1}{T} \sum_{i=1}^T e^{jz_i(X_i - X_j)} \\&= \frac{1}{T} \sum_{i=1}^T e^{jz_i X_i} e^{-jz_i X_j} \\&= \frac{1}{\sqrt{T}} \phi_z(X_i) \frac{1}{\sqrt{T}} \phi_z(X_j)\end{aligned}$$

$z_i \sim p(z)$ iid, Monte-Carlo, C

2. Fourier Random Features

- ▶ To compute $e^{iz_i X_i}$, use the cosine identity

$$e^{\pm j\theta} = \cos(\theta) \pm j\sin(\theta)$$

- ▶ $k(X_i - X_j)$ is a real-value of $e^{iz_i(X_i - X_j)}$
- ▶ The real-value of $e^{i(X_i - X_j)Z}$ is $\cos(X_i - X_j)Z$
- ▶ Two mapping of this exist as estimation:
 1. $\phi_z(X) = [\cos(XZ) \quad \sin(XZ)]$
 - ▶ This is a vector
 2. $\phi_z(X) = \sqrt{2} \cos(XZ + b)$
 - ▶ where b is drawn uniformly from $[0, 2\pi]$

Overall

- ▶ Let X be the dataset $\in R^{N \times D}$
- ▶ Let y be the target values $\in R^N$

1. Greedy Fitting

$$(W^*, Z^*) = \min_{W, Z} \left\| \sum_{i=1}^T w_i \phi_z(X_i) - y \right\|_F$$

2. Random Kitchen Sinks Fitting

$$z_1^*, z_2^*, \dots, z_T^* \sim p(w), \quad W^* = \min_W \left\| \sum_{i=1}^T w_i \phi_z(X_i) - y \right\|_F$$

1. Greedy method

- ▶ Inputs: y , probability measure μ
- ▶ Output $w_1, w_2, \dots, w_T, z_1, z_2, \dots, z_T$ so that

$$f(x) = \sum_{i=1}^{\infty} w_i \phi(x; z_i) \approx f_T(x) = \sum_{i=1}^T w_i \phi(x; z_i)$$

1. Initialize $f_0(x) = 0$
2. for $t = 1, 2, \dots, T$
 - ▶ $(z_t, w_t) = \arg \min_{w_t, z_t} \|(1 - w_t)f_{t-1} + \alpha_t \phi_{z_t}(x) - y\|_F$
 - ▶ $f_t \leftarrow f_{t-1} + \alpha_t \phi_{z_t}(X)$

2. Random Kitchen Sink method

- ▶ Inputs: y , probability measure p
- ▶ Output $w_1, w_2, \dots, w_T, z_1, z_2, \dots, z_T$ so that

$$f(x) = \sum_{i=1}^{\infty} w_i \phi(x; z_i) \approx f_T(x) = \sum_{i=1}^T w_i \phi(x; z_i)$$

1. Draw $z_1, z_2, \dots, z_T \sim p(z)$
2. $w \leftarrow \arg \min_w \|\sum_{i=1}^T w_i \phi_{z_i}(X) - y\|_F$

2. Random Kitchen Sink method in Python

```
# Random Kitchen Sink method

# Fit a Gaussian Process using kernel approximation
import numpy as np

# Stage 1. Get the random features
# 1. N samples and D features
N, D = 100, 50

# 2. Create Synthetic dataset
X = np.random.randn(N, D)
y = np.random.rand(N, 1)

# 3. Sample T from p(z)
T = 10
gamma = 1.
Z = (np.sqrt(2 * gamma) * np.random.randn(D, T))

# 4. Project X
phi = np.dot(X, Z)
b = np.random.uniform(0, 2 * np.pi, T)
phi += b
phi = np.cos(phi)
phi *= np.sqrt(2.) / np.sqrt(T)

# Stage 2. Fit ridge regression on the features
y_new = np.dot(phi.T, y)
lambda_ = 1. / D
W = (np.linalg.pinv(np.dot(phi.T, phi) +
                    np.identity(T) * lambda_)) * y_new
```

Fastfood

D dimension, T random features, N datapoints

- ▶ Sample $D \times T$ random numbers : $Z \sim N(0, \sigma^{-2})$
- ▶ $\phi_z(X_i) = e^{(-iX_iZ)}$
- ▶ Computing the features $\phi_z(\cdot)$ for every datapoint is $O(NDT)$
- ▶ Fastfood - using Hadamard transform we can reduce the time complexity for computing the features to $O(\log(D)NT)$
- ▶ For images, if $D = 10,000$, $\log_{10}(D) = 4$.

Fastfood

- ▶ Main idea: Compute $XZ \approx XV$
- ▶ V has similar properties to the Gaussian matrix Z

$$V = \frac{1}{\sigma\sqrt{N}} SHGP HB$$

- ▶ P is a $N \times N$ permutation matrix
- ▶ G is a diagonal random Gaussian
- ▶ B is a diagonal random $\{+1, -1\}$
- ▶ S is a diagonal random scaling
- ▶ H is Walsh- Hadamard matrix

More details

- ▶ Binary scaling matrix B :
 - ▶ It is a diagonal matrix with $B_{ii} \in \pm 1$ drawn iid.
- ▶ Permutation P :
 - ▶ can be generated by sorting random numbers.
- ▶ Gaussian scaling matrix G :
 - ▶ This is a diagonal matrix whose elements $G_{ii} \sim \mathcal{N}(0, 1)$ are drawn iid from a Gaussian.
- ▶ Scaling matrix S :
 - ▶ Gaussian case S ensures that the length distribution of the row of V are independent of each other.

More details

- ▶ Main idea: Compute $XZ \approx XV$
- ▶ V has similar properties to the Gaussian matrix Z

$$V = \frac{1}{\sigma\sqrt{N}}SHGPHB$$

- ▶ $S H G P H B$ produces pseudo-random Gaussian vectors
- ▶ S fixes the lengths to have the correct distribution

Analysis - Random Kitchen sink

- ▶ Computing Features $\phi(x) = \frac{1}{\sqrt{T}} \exp(iZX)$

- ▶ Regression:

$$w = [\phi(x)^T \phi(x)]^{-1} \phi(x)^T y$$

- ▶ Train time complexity

$$O(\underbrace{NTD}_{\text{Computing Random Features}} + \underbrace{T^3}_{\text{Inverting Covariance matrix}} + \underbrace{T^2 N}_{\text{Multiplication}})$$

- ▶ Prediction: $y^* = w^T \phi(x^*)$

- ▶ Test time complexity: $O(\underbrace{N^* TD}_{\text{Computing Random Features}} + \underbrace{T^2 N^*}_{\text{Multiplication}})$

Analysis - FastFood

- ▶ Computing Features $\phi(x) = \frac{1}{\sqrt{T}} \exp(iVx) \approx \frac{1}{\sqrt{T}} \exp(iZx)$
- ▶ regression:

$$w = [\phi(x)^T \phi(x)]^{-1} \phi(x)^T y$$

- ▶ FF Train time complexity

$$O(\underbrace{\log(N)TD}_{\text{Computing Random Features}} + \underbrace{T^3}_{\text{Inverting Covariance matrix}} + \underbrace{T^2N}_{\text{Multiplication}})$$

- ▶ RKS Train time complexity

$$O(\underbrace{NTD}_{\text{Computing Random Features}} + \underbrace{T^3}_{\text{Inverting Covariance matrix}} + \underbrace{T^2N}_{\text{Multiplication}})$$

- ▶ Prediction: $y^* = w^T \phi(x^*)$

- ▶ Test time complexity: $O(\underbrace{\log(D)TN^*}_{\text{Computing Random Features}} + \underbrace{T^2N^*}_{\text{Multiplication}})$

- ▶ For images, if $N = 10,000$, $\log_{10}(N) = 4$

Analysis - FastFood

We can lower the variance of the estimate of the kernel by concatenating D randomly chosen z_ω into one D -dimensional vector \mathbf{z} and normalizing each component by \sqrt{D} . The inner product $\mathbf{z}(\mathbf{x})'\mathbf{z}(\mathbf{y}) = \frac{1}{D} \sum_{j=1}^D z_{\omega_j}(\mathbf{x})z_{\omega_j}(\mathbf{y})$ is a sample average of z_ω and is therefore a lower variance approximation to the expectation (2).

Since z_ω is bounded between $+\sqrt{2}$ and $-\sqrt{2}$ for a *fixed* pair of points \mathbf{x} and \mathbf{y} , Hoeffding's inequality guarantees exponentially fast convergence in D between $\mathbf{z}(\mathbf{x})'\mathbf{z}(\mathbf{y})$ and $k(\mathbf{x}, \mathbf{y})$: $\Pr[|\mathbf{z}(\mathbf{x})'\mathbf{z}(\mathbf{y}) - k(\mathbf{x}, \mathbf{y})| \geq \epsilon] \leq 2 \exp(-D\epsilon^2/4)$. Building on this observation, a much stronger assertion can be proven for every pair of points in the input space simultaneously:

Figure 7:

Method	Train Time	Test Time	Train Mem	Test Mem
Naive	$\mathcal{O}(N^2 D)$	$\mathcal{O}(ND)$	$\mathcal{O}(ND)$	$\mathcal{O}(ND)$
Low Rank	$\mathcal{O}(NTD)$	$\mathcal{O}(TD)$	$\mathcal{O}(TD)$	$\mathcal{O}(TD)$
Kitchen Sinks	$\mathcal{O}(NTD)$	$\mathcal{O}(TD)$	$\mathcal{O}(TD)$	$\mathcal{O}(TD)$
Fastfood	$\mathcal{O}(NT \log(D))$	$\mathcal{O}(T \log(D))$	$\mathcal{O}(T \log(D))$	$\mathcal{O}(T)$

Figure 8:

Analysis - FastFood

- MNIST Dataset
 - 20000 training examples and 10000 test examples.

